

COMPSCI2030 Systems Programming

Sanitizers

Yehia Elkhatib



University
of Glasgow

Static Analysis

- A static analyzer reasons about the code *without* executing it
- The compiler performs some static analysis every time you compile code
 - e.g. type checking
- It is good practice to enable all warnings and make warnings errors
 - i.e. use the flags `-Wall -Werror`
 - so the compiler can be most helpful in detecting bugs before execution
- Some static analysis is too expensive to perform in every build
Other static analysis enforces a particular coding guideline
- We invoke the static analyzer using a flag and specifying the output format of the report:

```
clang --analyze --analyzer-output html program.c
```

clang static analyzer report

- The tool generates a report explaining potential bugs

Bug Summary

File: [REDACTED]/slides/SystemsProgramming/lecture5examples/clang-analyzer-examples/examples/4.c

Warning: [line 7, column 14](#)
The right operand of '<' is a garbage value

Annotated Source Code

Press ? to see keyboard shortcuts

Show analyzer invocation

☐ Show only relevant lines

```
1 // report two bugs about uninitialized value, currmin.
2
3 int minval(int *A, int n) {
4     int currmin, i;
5
6     while (i < n) {
7         if (A[i] < currmin)
8             currmin = A[i];
9         return currmin;
10 }
```

Here the tool has reported a “garbage value” which in this case leads to undefined outcome for an if-branch

clang-tidy

- A *linter* (the name comes from the first UNIX tool to perform static analysis on C)
- `clang-tidy` is a flexible tool that allows to enforce coding guidelines and to modernize source code
 - It is possible to extend `clang-tidy` by writing your own checks
- It is invoked like `clang`, accepting the same flags but after two dashes: `--`
- A series of *checks* can be enabled / disabled


```
$ clang-tidy -checks="readability-*" 6.c -- -Wall -Werror
/Users/lito/Desktop/examples/6.c:2:16: warning: pointer parameter 'p' can be pointer to const
[readability-non-const-parameter]
void test(int *p) {
      ~~~ ^
      const
/Users/lito/Desktop/examples/6.c:4:9: warning: statement should be inside braces
[readability-braces-around-statements]
if (p)
  ^
  {
```

Here we enable the checks for *readability*

It suggests to to use `const` and put braces around the branch of an if statement

- Detailed information is available at: <http://clang.llvm.org/extra/clang-tidy/>

Use the analyzer to improve your code



Dynamic Analysis Tools

- There exists a family of bug detection tools that use dynamic analysis
- These tools need the program to run and can only detect bugs which are encountered during the execution of a particular test input
- The `clang` project calls these tools *sanitizers*. The most important are:
 - `AddressSanitizer` - a memory error detector
 - `MemorySanitizer` - a detector of uninitialized reads
 - `LeakSanitizer` - a memory leak detector
 - `UndefinedBehaviorSanitizer` - a detector of undefined behaviour
- Later in the course, you might want to look up:
 - `ThreadSanitizer` - a data race detector

Address Sanitizer

- Address Sanitizer is a memory error detector for:
 - Out-of-bounds / Use-after-free / Double free memory accesses
- Makes clang insert instructions to monitor every memory access
- This slows down the execution by about 2x
 - valgrind (a similar tool) has often a slowdown of 20-100x!
- These flags enable address sanitizer:

```
clang -fsanitize=address -fno-omit-frame-pointer -O1 -g -Wall -Werror program.c -o program
```

 - `fno-omit-frame-pointer` produces a readable call stack
 - `O1` enables basic optimizations
- The compiler will produce a binary as usual: `./program`
- Address Sanitizer has found hundreds of bugs in large scale software projects
 - e.g. Chromium and Firefox

Address Sanitizer output

```
seAfterFree

22:55:53 in ↗slides/SystemsProgramming on ↗ master [?]
λ → ./useAfterFree
=====
==6954==ERROR: AddressSanitizer: heap-use-after-free on address 0x614000000044 at pc 0x000108f72ef8 bp 0x7ffee6c8d3d0 sp 0x7ffee6c8d3c8
READ of size 4 at 0x614000000044 thread T0
    #0 0x108f72ef7 in main useAfterFree.c:6
    #1 0x7fff76bdc084 in start (libdyld.dylib:x86_64+0x17084)

0x614000000044 is located 4 bytes inside of 400-byte region [0x614000000040,0x6140000001d0)
freed by thread T0 here:
    #0 0x108fce10d in wrap_free (libclang_rt.asan_osx_dynamic.dylib:x86_64h+0x5710d)
    #1 0x108f72ebe in main useAfterFree.c:5
    #2 0x7fff76bdc084 in start (libdyld.dylib:x86_64+0x17084)

previously allocated by thread T0 here:
    #0 0x108fcdf53 in wrap_malloc (libclang_rt.asan_osx_dynamic.dylib:x86_64h+0x56f53)
    #1 0x108f72eb3 in main useAfterFree.c:4
    #2 0x7fff76bdc084 in start (libdyld.dylib:x86_64+0x17084)

SUMMARY: AddressSanitizer: heap-use-after-free useAfterFree.c:6 in main
Shadow bytes around the buggy address:
  0x1c27fffffffb0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  0x1c27fffffffc0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  0x1c27ffffffd0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  0x1c27ffffffe0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
  0x1c27fffffff0: 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
⇒0x1c2800000000: fa fa fa fa fa fa fa fa fd fd fd fd fd fd fd fd
  0x1c2800000010: fd fd fd fd fd fd fd fd fd fd fd fd fd fd fd fd
  0x1c2800000020: fd fd fd fd fd fd fd fd fd fd fd fd fd fd fd fd
  0x1c2800000030: fd fd fd fd fd fd fd fd fd fd fd fa fa fa fa fa fa
  0x1c2800000040: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
  0x1c2800000050: fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa fa
Shadow byte legend (one shadow byte represents 8 application bytes):
Addressable: 00
Partially addressable: 01 02 03 04 05 06 07
Not addressable: 0f 10 11 12 13 14 15 16 17 18 19 1a 1b 1c 1d 1e 1f
Hugepage not available: 20 21 22 23 24 25 26 27
==6954==ABORTING
```

Reports a heap-use-after-free on address 0x614000000044

Provides information where the memory was freed (line 5) and allocated (line 4)

Memory Sanitizer

- Detects uninitialized reads to memory

```
#include <stdlib.h>
#include <stdio.h>
int main(int argc, char** argv) {
    int* a = malloc(sizeof(int)*10);
    a[5] = 0;
    if (a[argc])
        printf("xx\n");
    return 0;
}

% clang -fsanitize=memory -fno-omit-frame-pointer -g -O2 umr.cc
% ./a.out
WARNING: MemorySanitizer: use-of-uninitialized-value
#0 0x7f45944b418a in main umr.c:6
#1 0x7f45938b676c in __libc_start_main libc-start.c:226
```

- Under active development
- Currently only available for Linux and BSD

Leak Sanitizer

- Leak sanitizer detects memory leaks (i.e. un-free'd memory blocks)

```
[~bash-4.2$ cat mem_leak.c
#include <stdlib.h>

int main() {
    void *p = malloc(7);
    p = 0; // The memory is leaked here.
    return 0;
}

[~bash-4.2$ clang -fsanitize=address -g mem_leak.c
[~bash-4.2$ ASAN_OPTIONS=detect_leaks=1 ./a.out

=====
==83220==ERROR: LeakSanitizer: detected memory leaks

Direct leak of 7 byte(s) in 1 object(s) allocated from:
    #0 0x465269 in __interceptor_malloc (/users/staff/yehia/sp-ga/a.out+0x465269)
    #1 0x47b54c in main /users/staff/yehia/sp-ga/mem_leak.c:4
    #2 0x7f4631ba3554 in __libc_start_main (/lib64/libc.so.6+0x22554)

SUMMARY: AddressSanitizer: 7 byte(s) leaked in 1 allocation(s).
```

- Under active development
- Available for Linux, macOS, NetBSD

Undefined Behaviour

- For certain operations, the C standard demands no particular behaviour
 - typically cases that are considered bugs, e.g. dereferencing a `null` pointer
- It is expensive to check if dereferencing a pointer is valid every time
- Because the compiler does not have to ensure a particular behaviour for `null` pointers, it *assumes* that the programmer checked that the pointer is not `null`
- Undefined behaviour is therefore crucial for fast code, but makes detection of bugs much harder, as it is not guaranteed that a bug will result in a crash
- A good introduction to undefined behaviour is this series of blog posts: [What Every C Programmer Should Know About Undefined Behavior](#)

Undefined Behaviour Sanitizer

- Detects various types of undefined behaviour

```
int main(int argc, char **argv) {  
    int k = 0x7fffffff; // this is the largest possible signed int value ...  
    k += argc; // ... this will produce an integer overflow  
    return 0;  
}
```

```
% clang -fsanitize=undefined -Wall -Werror intOverflow.c -o intOverflow  
% ./intOverflow  
intOverflow.c:3:5: runtime error: signed integer overflow: 2147483647 + 1 cannot  
be represented in type 'int'
```



An integer overflow is
detected