

COMPSCI2030 Systems Programming

Dynamic Memory Allocation


Yehia Elkhatib



University
of Glasgow

Stack vs. Heap memory regions

- So far we have only seen variables with automatic allocation
 - lifetime managed by the compiler
- These variables are stored in the stack part of the memory
 - the size of every variable on the stack has to be known *statically*, i.e. without executing the program
- In many cases we do not know the size of a variable statically
 - e.g. dynamically sized arrays
- For such cases we manage the memory manually by dynamically *requesting* and *freeing* memory
- The part of memory used for dynamic allocation is the heap
- The stack and the heap share the same address space and grow with use towards each other



stack and heap are in a single address space

Dynamic memory management

All defined in
<stdlib.h>

- We request a new chunk of memory from the heap

```
void *malloc( size_t size );
```
- We specify the number of bytes we like to allocate
 - success: it returns a void-pointer to the first byte of the un-initialised block
 - failure: it returns NULL; so it is good practice to check directly afterwards
- Other related functions
 - Reallocate a given area of memory

```
void* realloc( void *ptr, size_t new_size );
```
 - Allocate an array of num objects of size, initialize all to zero

```
void* calloc( size_t num, size_t size );
```

Dynamic memory management

- Memory allocated with malloc, etc. must be manually deallocated
`void free(void *ptr);`
- If free is not called, we *leak* the allocated memory
 - i.e. memory blocks are allocated but not used
- Good practice: set a pointer to NULL after freeing it
 - makes it easy to check if you have free'd it or not



Example: Dynamically sized array

- The size of the array is the first number entered by the user

```
#include <stdlib.h>
#include <stdio.h>
int main() {
    printf("How many numbers do you want to average?\n");
    int count;
    if (scanf("%d", &count) == EOF) { exit(-1); }
    // allocate memory based on dynamic input (here its size)
    int* array = malloc(count * sizeof(int));
    for (int i = 0; i < count; i++) {
        int number;
        if (scanf("%d", &number) == EOF) { exit(-1); }
        array[i] = number;
    }
    float sum = 0.0f;
    for (int i = 0; i < count; i++) { sum += array[i]; }
    printf("The average is %.2f\n", sum / count);
    free(array); array=NULL; // free the memory manually after use
}
```

Returning a pointer to a local variable

- It is an easy mistake to return a pointer to a local variable. Never do it!
- Because the pointer has a longer lifetime than the variable it is/was pointing to

```
struct node* create_node(char value) {  
    struct node node;  
    node.value = value;  
    node.next = NULL;  
    return &node;  
} // lifetime of node ends here...  
    // but its address lives on in a_ptr  
  
int main() {  
    struct node* a_ptr = create_node('a');  
    // ...  
} // lifetime of a_ptr (pointing to node) ends here
```

- Solutions

- Pass the root/leading node, then manipulate it inside the function before node expires
- Allocate the memory and pass the pointer back

Function Pointers

- Memory does not store data but also program code
- It is possible to have a pointer pointing to code
- These pointers are called *function pointers* and have the type:
 `return_type (*) (argument_types)`
- This is typically used when passing to another function
 - i.e. telling the second function which function to use *at runtime*
 - very useful with design patterns such as command, iterator, visitor
- Example:
 - a sorting function that should compare elements according to their type
 - numbers >; strings `strcmp()`; ...

Function Pointers

- Now we can implement a generic print function for a binary tree:

```
void print_string(const void * value_ptr) {  
    char * string = value_ptr; // by changing the type we give the bits meaning  
    printf("%s\n", string);  
}  
  
void print(node * root, void (* print_function)(const void *) ) {  
    if (root) {  
        print_function(root->value_ptr);  
        print(root->left_child, print_function);  
        print(root->right_child, print_function);  
    }  
}  
  
int main() {  
    node * root = ... ;  
    print(root, print_string);  
}
```

- Function names are automatically converted to function pointers
 - i.e. we do not have to write &print_string

Memory Management Challenges

- When we allocate memory with `malloc`, we are responsible for calling `free`
- We must call `free` exactly once for each `malloc`'d address
- Good practice: assign `NULL` value to pointers that have been free'd
- But this does *not* prevent all double free errors, e.g.:

```
int main() {  
    void * ptr1 = malloc(sizeof(int));  
    void * ptr2 = ptr1;  
    free(ptr1); ptr1 = NULL;  
    free(ptr1); // OK. Calling free with NULL is fine  
    free(ptr2); // *** error for object 0x7f9355c00690: pointer being freed  
                // was not allocated  
}
```

Memory Management Challenges

- Another problem are dangling pointers: these point to locations that have been free'd

```
int main() {  
    node * left_child = create_tree("a", NULL, NULL);  
    node * root = create_tree("b", left_child, NULL );  
    destroy_tree(left_child); // now: root->left_child points to freed memory!  
}
```

Memory Management Challenges

- Inconspicuous memory leaks

```
#include <stdlib.h>
#include <stdio.h>

int main() {
    char *mem = (char*) malloc( sizeof(char) * 20); // allocate some memory
    if (!mem) { exit(EXIT_FAILURE); } // check if allocation went fine
    // use allocated memory

    mem = (char*) malloc( sizeof(char) * 30); // allocate more memory
    // ...BUT we just lost the pointer to the old memory => LEAK!

    // use newly allocated memory

    free(mem); // free newly allocated memory
}
```