

Systems Programming – Part 2

Concurrent Systems Programming

Dr Lauritz Thamsen
lauritz.thamsen@glasgow.ac.uk
<https://lauritzthamsen.org>

- **Intro to Concurrency (with Processes and Threads)**
 - Intro to POSIX Threads
- Process/Thread Synchronisation
- More on Process Management (from an OS Perspective)
- Concurrency Beyond Threads & Limits of Scalability
- Virtual Memory & Levels of Storage

Intro to POSIX Threads

Thread Implementations

- There are thread implementations in almost all programming languages; we will look at the C **pthread** library
- Many threading implementations are conceptually quite similar, and realise a particular thread lifecycle:
 - A thread is **created**: starts executing a specified function, with some arguments, and is given an identifier
 - A thread can **wait** for another thread to terminate
 - A thread can **cancel** another thread
 - A thread **terminates** either by calling exit or when its main function ends
 - **Communication** between threads happens through modifying the state of **shared variables**

- Part of the **POSIX specification** collection, defining an API for thread creation and management (*pthread.h*)
- Implementations for all Unix-alike operating systems available
 - Utilization of kernel- or user-mode threads depends on implementation
- Groups of functionality (pthread_ function prefix)
 - **Thread management**: Start, wait for termination, ...
 - **Synchronization means**: Mutexes, condition variables, read/write locks, barriers
- Semaphore API is a separate POSIX specification (sem_ prefix)

POSIX Threads

- POSIX Threads is the most used threading implementation for C
- To use it we need to *#include* `<pthread.h>` and specify a compiler flag (`-lpthread`)

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>
#include <assert.h>

void * PrintHelloWorld(void*) {
    printf("Hello World from a thread!\n");
    return NULL;
}
```

```
int main() {
    pthread_t thread;
    int error = pthread_create(&thread,
        NULL, PrintHelloWorld, NULL);
    assert(error == 0);
    printf("Created thread\n");
    error = pthread_join(thread, NULL);
    assert(error == 0);
}
```

```
clang -Wall -Werror program.c -lpthread -o program
```

- *pthread_create()*
 - Create a new thread in the same process, with a given routine and argument
- *pthread_exit()*, *pthread_cancel()*
 - Terminate a thread from inside (→ *exit*) or outside (→ *cancel*) of the thread

Creating POSIX Threads

- Threads are created with the `pthread_create` function

```
int pthread_create(pthread_t *thread,  
                  const pthread_attr_t *attr,  
                  void *(*start_routine)(void*),  
                  void *arg);
```

- It takes four arguments:
 1. A **thread** identifier, which is a pointer to a memory location of type `pthread_t`
 2. Thread attributes which set properties such as scheduling policies or stack size (and passing `NULL` results in default attributes)
 3. A **function pointer** to the ***start_routine***
This function takes a single argument of type `void*` and returns a value of type `void*`
 4. The **argument** that is passed **to *start_routine*** (a pointer)
- It returns 0 if the thread is created successfully or, else, a non-zero error code
- **Passing pointers to and from `start_routine` allows the passing of arbitrary data**
- It requires care to ensure that the memory locations pointed to have appropriate lifetimes

Passing an argument to a thread

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#define NUM_THREADS 5

void *PrintHello(void *threadid)
{
    long tid;
    tid = (long)threadid;
    printf("Hello World! It's me, thread #%ld!\n", tid);
    pthread_exit(NULL);
}

int main(int argc, char *argv[])
{
    pthread_t threads[NUM_THREADS];
    int rc;
    long t;
    for(t=0;t<NUM_THREADS;t++){
        printf("In main: creating thread %ld\n", t);
        rc = pthread_create(&threads[t], NULL, PrintHello, (void *)t);
        if (rc){
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

Waiting for Other Threads

- *pthread_join()*
 - Blocks the caller until the specified thread terminates
 - Only one joining thread per target is allowed
 - If the thread gave an exit code to *pthread_exit()*, it can be determined here

Waiting for Other Threads

- To wait for another thread to terminate we use *pthread_join*

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- It takes two arguments:
 1. A thread identifier (the thread we want to wait for)
 2. A **pointer** to a **memory location** of type `void*`

The return value of the *start_routine* (== a generic pointer), as passed to *pthread_create*, will be copied to this location
- It returns 0 on success and, otherwise, a non-zero error code

Waiting for Another Thread

- Example of returning a single int value from a thread:

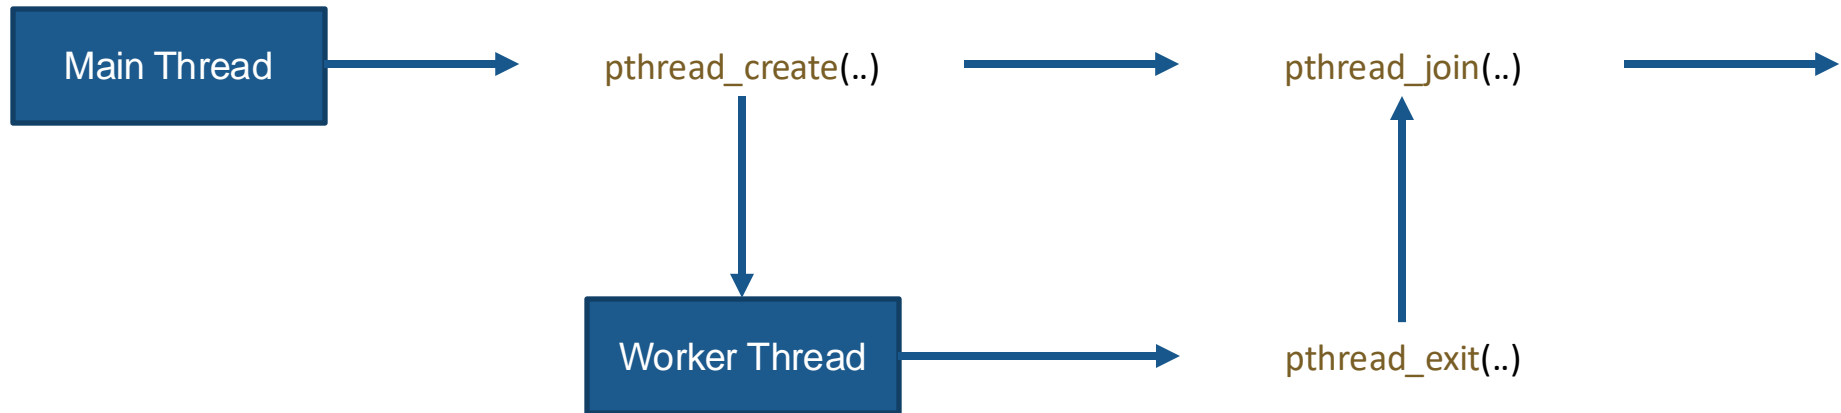
```
int* return_value;  
  
int error = pthread_join(thread, (void**)&return_value);  
  
assert(error == 0);  
  
if (return_value) {  
    printf("return_value: \n", *return_value);  
}  
  
// maybe: free(return_value);
```

Waiting for Another Thread

```
void *thread_func(void *args) {  
    /* ... */  
  
    /* Allocate heap space for this thread's results */  
    struct results *results = calloc(sizeof (struct results), 1);  
  
    /* ... populate results struct ... */  
  
    /* Return the pointer to results */  
    pthread_exit(results);  
}
```

```
struct results *results;  
  
int error = pthread_join(thread, (void **)&results);  
  
assert(error == 0);  
  
/* make use of results, but then also free(results)  
at some point */
```

Forking and Joining Threads



Completing this activity will support your understanding of the material today (before we have real lab time tomorrow!)

In the “Lecture 2.2 files” folder on Moodle

- Review the code for *pthread_hello_world.c*
- Compile it and run it – don’t forget the “-lpthread” flag for clang (see Slide 5 for an example)
- Add another print statement “Last thread about to die” just before the main thread terminates

(Optional) Recommended Reading

- Blaise Barney, Lawrence Livermore National Laboratory, UCRL-MI-133316, POSIX Threads Programming, <https://hpc-tutorials.llnl.gov/posix/>, Section 5 for today and later Section 7
- Videos
 - Programming with Threads, Jacob Sorber, <https://www.youtube.com/playlist?list=PL9IEJIKnBJjFZxuqyJ9JqVYmuFZHr7CFM>
 - Unix Threads in C, CodeVault, <https://www.youtube.com/playlist?list=PLfqABt5AS4FmuQf70psXrsMLEDQXNkLq2>